

# A Dual-SPMA Framework for Convex MDPs

## Fenchel Duality + Softmax Policy Mirror Ascent

Shervin Khamooshian   Ahmed Magd   Pegah Aryadoost   Danielle Nguyen

School of Computing Science, Simon Fraser University

Project Presentation

### Main Claim

Fenchel duality + a fast policy optimizer (SPMA) gives a simple, competitive way to solve convex MDPs; we compare this Dual-SPMA recipe against NPG-PD.

# Outline


- ① Background: Reinforcement Learning
- ② Motivation: Convex MDPs
- ③ Problem Formulation: Fenchel Duality
- ④ Related Work: “Reward Is Enough” & SPMA
- ⑤ Our Method: Dual-SPMA
- ⑥ Experiments
- ⑦ Conclusion & Future Work

# What is Reinforcement Learning?

**Reinforcement Learning (RL)** is a learning framework where an agent learns to make decisions by interacting with an environment.

At each time step  $t$ , the agent:

- ① Observes a **state**  $s_t$
- ② Chooses an **action**  $a_t$  (based on a policy)
- ③ Receives a **reward**  $r_t$
- ④ Transitions to a new **state**  $s_{t+1}$



# Markov Decision Process (MDP): Formal Definition

An **MDP** is defined by the tuple  $(\mathcal{S}, \mathcal{A}, P, r, \gamma, \rho)$ :

---


Symbol	Meaning
$\mathcal{S}$	State space (set of all possible states)
$\mathcal{A}$	Action space (set of all possible actions)
$P(s' s, a)$	Transition probability: probability of reaching $s'$ from $(s, a)$
$r(s, a)$	Reward function: immediate reward for taking action $a$ in state $s$
$\gamma \in [0, 1]$	Discount factor: how much to value future vs. immediate rewards
$\rho(s)$	Initial state distribution

---

**Policy**  $\pi(a|s)$ : probability of taking action  $a$  in state  $s$ .

# Trajectory and Return

A **trajectory**  $\tau$  is a sequence of states, actions, and rewards:



## Discounted Return

The **expected discounted return** under policy  $\pi$  is:  $J(\pi) = \mathbb{E}_\pi [\sum_{t=0}^{\infty} \gamma^t r(s_t, a_t)]$

**Goal of RL:** Find  $\pi^* = \arg \max_\pi J(\pi)$


# Occupancy Measure: Where the Policy Spends Time

## Discounted Occupancy Measure

$$d_{\pi}(s, a) = (1 - \gamma) \sum_{t=0}^{\infty} \gamma^t \Pr_{\pi}(s_t = s, a_t = a)$$

### Notation:

- $d_{\pi}(s, a)$ : probability of being in state  $s$  and taking action  $a$  under policy  $\pi$
- $(1 - \gamma)$ : normalization factor
- $\Pr_{\pi}$ : probability under policy  $\pi$



Think of  $d_{\pi}$  as a **heatmap**: bright = often visited

**Key property:**  $d_{\pi}$  is a probability distribution:  $\sum_{s,a} d_{\pi}(s, a) = 1$

# Key Identity: RL is Linear in Occupancy

## Fundamental Identity

$$\langle r, d_\pi \rangle = \sum_{s,a} r(s, a) d_\pi(s, a) = (1 - \gamma) J(\pi)$$

## What does this mean?

- Up to the constant  $(1 - \gamma)$ , the RL objective is **linear** in  $d_\pi$
- Weights = how often we visit each  $(s, a)$  pair
- Standard RL  $\Rightarrow$  maximize a **linear** function of  $d_\pi$

But many interesting goals are NOT linear in  $d_\pi$ :

Goal	Objective
Safety constraints	$\max J_r(\pi)$ subject to $\langle c, d_\pi \rangle \leq \tau$
Imitation learning	$\min \ d_\pi - d_{\text{expert}}\ $
Exploration	$\max J_r(\pi) + \alpha H(d_\pi)$

# Why Convex MDPs?

**Problem:** Linear RL is insufficient for:

- Safety constraints
- Matching expert behavior
- Encouraging exploration
- Risk-sensitive objectives

**Solution:** Convex MDPs

$$\min_{\pi} f(d_{\pi})$$

where  $f$  is a **convex function**

**Challenge:**

- Optimizing over occupancy measures is hard
- High-dimensional constrained space
- Can't directly apply standard RL

**Our approach:**

- Use **Fenchel duality**
- Transform to min-max game
- Reduce to shaped-reward RL

# Convex MDP: Examples

## General Form

$\min_{d \in \mathcal{D}} f(d)$  where  $\mathcal{D}$  = feasible occupancy measures,  $f$  = convex

### Example 1: Standard RL (linear, trivially convex)

$$f(d) = -\langle r, d \rangle = -\sum_{s,a} r(s, a) d(s, a)$$

### Example 2: Entropy-Regularized RL

$$f(d) = -\langle r, d \rangle + \alpha \sum_{s,a} d(s, a) \log d(s, a)$$

### Example 3: Constrained Safety (CMDP)

$$f(d) = -\langle r, d \rangle + \mu \max\{0, \langle c, d \rangle - \tau\}$$

where  $c(s, a)$  = cost function,  $\tau$  = threshold,  $\mu$  = fixed penalty weight

# Roadmap

- ✓ Background: Reinforcement Learning

- ✓ Motivation: Convex MDPs

- **Problem Formulation: Fenchel Duality**

- Related Work: “Reward Is Enough” & SPMA

- Our Method: Dual-SPMA

- Experiments

- Conclusion & Future Work

# Fenchel Conjugate: Definition

Given a convex function  $f : \mathbb{R}^n \rightarrow \mathbb{R} \cup \{\infty\}$

## Fenchel Conjugate (Convex Conjugate)

$$f^*(y) = \sup_{x \in \mathbb{R}^n} \{\langle y, x \rangle - f(x)\}$$

### Notation:

- $f^*(y)$ : the conjugate function evaluated at dual variable  $y$
- $\langle y, x \rangle = \sum_i y_i x_i$ : inner product (dot product)
- $\sup$ : supremum (least upper bound)

**Intuition:**  $f^*(y)$  measures “how much  $\langle y, x \rangle$  can exceed  $f(x)$ ”

# Fenchel–Moreau Theorem

## Fenchel–Moreau Identity

For any proper, closed, convex function  $f$ :

$$f(d) = \sup_y \{ \langle y, d \rangle - f^*(y) \}$$

### What does this say?

- We can **recover**  $f$  from its conjugate  $f^*$
- $f$  is the conjugate of its conjugate:  $f = (f^*)^*$
- This is called **biconjugation**

### Why is this useful?

- Transforms a minimization problem into a **min-max** problem
- Introduces a **dual variable**  $y$  that we can optimize over

# Applying Fenchel Duality to Convex MDPs

**Step 1:** Start with the convex MDP problem

$$\min_{d \in \mathcal{D}} f(d)$$

**Step 2:** Apply Fenchel–Moreau identity

$$\min_{d \in \mathcal{D}} f(d) = \min_{d \in \mathcal{D}} \sup_y \{ \langle y, d \rangle - f^*(y) \}$$

**Step 3:** This is a **convex-concave saddle-point problem**

$$= \min_{d \in \mathcal{D}} \max_y \{ \langle y, d \rangle - f^*(y) \}$$

(Under standard conditions, solving this saddle-point is equivalent to the original problem.)

**Step 4:** Replace  $d$  with  $d_\pi$  (occupancy induced by policy)

## Saddle-Point Formulation

$$\min_{\pi} \max_y \underbrace{\langle y, d_\pi \rangle - f^*(y)}_{L(\pi, y)}$$

# Two-Player Game Interpretation

## Saddle-Point Problem

$$\min_{\pi} \max_y L(\pi, y), \quad \text{where } L(\pi, y) = \langle y, d_{\pi} \rangle - f^*(y)$$

This is a **min-max game** between two players:

### Policy Player (min)

- Chooses policy  $\pi$
- Wants to minimize  $L$
- Controls occupancy  $d_{\pi}$

### Dual Player (max)

- Chooses dual variable  $y$
- Wants to maximize  $L$
- Shapes the reward signal

At equilibrium: policy player finds optimal  $\pi^*$ ,  
dual player finds optimal  $y^*$

## From Saddle Point to Shaped Reward

For **fixed** dual variable  $y$ , the policy player solves:  $\min_{\pi} \langle y, d_{\pi} \rangle$

Expand using the occupancy definition:

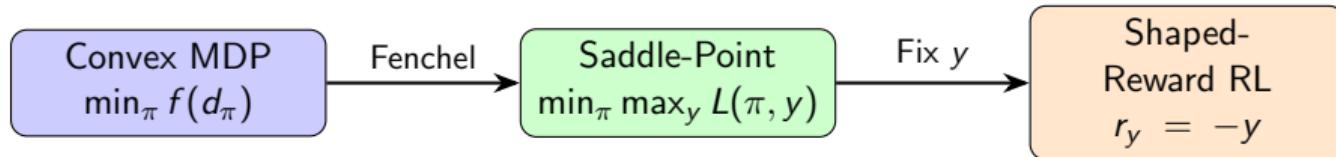
$$\langle y, d_{\pi} \rangle = \sum_{s,a} y(s, a) \cdot d_{\pi}(s, a) = (1 - \gamma) \mathbb{E}_{\pi} \left[ \sum_{t=0}^{\infty} \gamma^t y(s_t, a_t) \right]$$

### Key Insight: Shaped Reward

$$\min_{\pi} \langle y, d_{\pi} \rangle = \min_{\pi} \mathbb{E}_{\pi} \left[ \sum_{t=0}^{\infty} \gamma^t y(s_t, a_t) \right] = \max_{\pi} \mathbb{E}_{\pi} \left[ \sum_{t=0}^{\infty} \gamma^t \underbrace{(-y(s_t, a_t))}_{r_y(s_t, a_t)} \right]$$

**Conclusion:** Policy player just does **standard RL** with shaped reward:  $r_y(s, a) = -y(s, a)$

# Summary: The Fenchel Dual Reduction



## Algorithm structure:

- ➊ **Dual step:** Update  $y$  using gradient of  $L$  w.r.t.  $y$

$$y_{k+1} = y_k + \alpha (d_{\pi_k} - \nabla f^*(y_k))$$

- ➋ **Policy step:** Run RL algorithm with reward  $r_{y_k} = -y_k$

$$\pi_{k+1} = \text{RL-Update}(\pi_k, r_{y_k})$$

## Related Work: “Reward Is Enough” (Zahavy et al., 2021)

### Main contributions of this foundational paper:

#### ① Fenchel dual reduction:

- Reformulate convex MDP as saddle-point problem
- The theoretical foundation we just presented

#### ② Meta-algorithm:

- Alternating updates between policy and dual players
- Any RL algorithm can be the policy player
- Any online convex optimization (OCO) can be the dual player

#### ③ Unification:

- Shows many RL paradigms are special cases of convex MDPs
- Imitation learning, constrained RL, entropy-regularized RL

**Our contribution:** Implement this framework with SPMA as the policy player

## Related Work: Softmax Policy Mirror Ascent (Asad et al., 2024)

### Softmax Policy Mirror Ascent (SPMA):

- Mirror ascent in *logit space* using log-sum-exp mirror map
- Achieves **linear convergence** in tabular MDPs

#### Tabular SPMA Update Rule

$$\pi_{t+1}(a|s) = \pi_t(a|s) \cdot (1 + \eta A^{\pi_t}(s, a))$$

where  $A^\pi(s, a) = Q^\pi(s, a) - V^\pi(s)$  is the advantage function<sup>1</sup>

#### Why use SPMA as our policy player?

- **Geometry-aware:** Updates respect the simplex structure
- **No normalization:** Unlike vanilla PG, no per-state renormalization
- **Fast:** Linear convergence vs. sublinear for vanilla PG
- **Function approximation:** Clean extension via convex classification

<sup>1</sup>In tabular theory, this uses  $[1 + \eta A]_+$  and implicit renormalization; simplified form shown here.

## Related Work: NPG–PD (Our Baseline)

**Natural Policy Gradient Primal–Dual** (Ding et al., 2020):

For constrained MDPs with Lagrangian:

CMDP Lagrangian

$$L(\pi, \lambda) = J_r(\pi) + \lambda(J_c(\pi) - \tau), \quad \lambda \geq 0$$

**Dual step (constraint):**

**Primal step (policy):**

- Natural policy gradient ascent
- Uses Fisher information matrix
- Geometry-aware like SPMA

$$\lambda_{k+1} = [\lambda_k + \beta(J_c(\pi_k) - \tau)]_+$$

- Projected gradient ascent
- $[\cdot]_+ = \max(0, \cdot)$

**Guarantees:**  $\mathcal{O}(1/\sqrt{T})$  optimality gap and constraint violation

# Roadmap

- ✓ Background: Reinforcement Learning
- ✓ Motivation: Convex MDPs
- ✓ Problem Formulation: Fenchel Duality
- ✓ Related Work: “Reward Is Enough” & SPMA

→ **Our Method: Dual–SPMA**

Experiments

Conclusion & Future Work

# Our Contributions

## ① Dual-SPMA Framework:

- Complete implementation of outer dual loop + SPMA policy oracle
- Supports entropy-regularized RL and constrained safety (CMDP)

## ② Three Occupancy Estimators:

- Tabular Monte Carlo
- Feature-based Monte Carlo
- MLE-style estimator (following Barakat et al., 2024)

## ③ NPG-PD Baseline:

- Faithful implementation for fair comparison
- Same architecture and hyperparameters where possible

## ④ Empirical Comparison:

- SPMA vs NPG-PD on constrained safety tasks

# Dual-SPMA Loop: High-Level View

## Saddle-Point Problem


$$\min_{\pi} \max_y \underbrace{\langle y, d_{\pi} \rangle - f^*(y)}_{L(\pi, y)}$$

Outer loop (dual):

$$y_{k+1} = y_k + \alpha(\hat{d}_{\pi_k} - \nabla f^*(y_k))$$

Inner loop (policy):

- Run  $K_{\text{in}}$  SPMA steps
- Shaped reward:  $r_{y_k} = -y_k$



## Policy Player: SPMA Actor Loss

**Standard Policy Gradient Loss:**

$$\mathcal{L}_{\text{PG}} = -\mathbb{E} [\log \pi(a|s) \cdot A(s, a)]$$

**SPMA adds a “stay close” regularizer:**

SPMA Actor Loss

$$\mathcal{L}_{\text{SPMA}} = \mathbb{E} \left[ -\Delta \log \pi \cdot A + \frac{1}{\eta} \underbrace{(\exp(\Delta \log \pi) - 1 - \Delta \log \pi)}_{\text{KL-like regularizer}} \right]$$

**Notation:**

- $\Delta \log \pi = \log \pi_{\text{new}}(a|s) - \log \pi_{\text{old}}(a|s)$ : change in log-probability
- $A(s, a)$ : advantage function
- $\eta$ : step size parameter (chosen via Armijo line search)

# Occupancy Estimation: MC vs MLE

## Monte Carlo (our default)

### Tabular Estimator

$$\hat{d}_\pi(s, a) = \frac{1-\gamma}{N} \sum_{i=1}^N \sum_{t=0}^T \gamma^t \mathbf{1}\{s_t^{(i)} = s, a_t^{(i)} = a\}$$

- Simple: count discounted visits
- Variance grows with  $|\mathcal{S}||\mathcal{A}|$

## MLE (Barakat et al., 2024)

### Log-Linear Model

$$\lambda_\omega(s, a) \propto \exp(\omega^\top \phi(s, a))$$

- Fit  $\omega$  by max-likelihood
- Error:  $O(\sqrt{m/n})$
- Independent of  $|\mathcal{S}||\mathcal{A}|$ !

**Sanity check:**  $\sum_{s,a} \hat{d}_\pi(s, a) \approx 1$  verified in all our tests

## Baseline: NPG-PD Implementation

Same Lagrangian as Dual-SPMA:

$$L(\pi, \lambda) = J_r(\pi) + \lambda(J_c(\pi) - \tau)$$

Primal (policy):

- Natural PG on shaped reward  
 $r_\lambda = r - \lambda c$
- Diagonal Fisher approximation
- Same actor-critic as SPMA

Dual (constraint):

$$\lambda_{k+1} = [\lambda_k + \beta(J_c - \tau)]_+$$

- One NPG step per iteration
- Evaluate  $J_c$  via Monte Carlo

Fair comparison: Same networks, same hyperparameters where possible

## Example: Constrained Safety CMDP

**Problem:** Maximize reward subject to safety constraint

$$\max_{\pi} J_r(\pi) \quad \text{s.t.} \quad J_c(\pi) \leq \tau$$

where  $J_r(\pi) = \mathbb{E}_{\pi}[\sum_t \gamma^t r(s_t, a_t)]$  and  $J_c(\pi) = \mathbb{E}_{\pi}[\sum_t \gamma^t c(s_t, a_t)]$ .

**Dual-SPMA approach:**

- ① Build dual variable:  $y_\lambda(s, a) = \lambda c(s, a) - r(s, a)$
- ② Policy sees shaped reward:  $r_y = -y = r - \lambda c$
- ③ Run SPMA inner loop on  $r_y$
- ④ Update dual:  $\lambda_{k+1} = [\lambda_k + \beta(J_c(\pi_k) - \tau)]_+$

**SPMA vs NPG-PD:** Same dual update, different policy optimizer!

Only difference is Step 3: SPMA loss vs. natural gradient

# Roadmap

- ✓ Background: Reinforcement Learning
- ✓ Motivation: Convex MDPs
- ✓ Problem Formulation: Fenchel Duality
- ✓ Related Work: “Reward Is Enough” & SPMA
- ✓ Our Method: Dual-SPMA

→ **Experiments**

Conclusion & Future Work

# Experimental Setup

## Environments:

- FrozenLake 4×4 (tabular)
- Deterministic transitions
- Unsafe states = holes (cost  $c = 1$ )

## Methods Compared:

- Dual-SPMA (ours)
- NPG-PD baseline

## Metrics:

- $J_r(\pi)$ : reward return
- $J_c(\pi)$ : cost return
- $J_c - \tau$ : constraint violation
- $\sum \hat{d}_\pi$ : estimator sanity

## Hyperparameters:

- Discount  $\gamma = 0.99$
- Safety threshold  $\tau = 0.1$
- 30 outer iterations
- 2048 steps/rollout

## Results: Entropy-Regularized RL

$L(\pi, y)$  vs iterations

(Add plot here)

Saddle value  $L(\pi, y)$  vs iterations

$\sum \hat{d}_\pi$  vs iterations

(Add plot here)

$\sum_{s,a} \hat{d}_\pi(s, a)$  vs iterations

### Observations:

- Saddle value  $L(\pi, y)$  converges smoothly
- Occupancy estimate stays near 1 (estimator is consistent)

## Results: Constrained Safety (Dual-SPMA)

$J_r, J_c$  vs iterations

(Add plot here)

$J_r(\pi_k)$  and  $J_c(\pi_k)$  vs iterations

Constraint violation

(Add plot here)

Constraint violation  $J_c - \tau$  and  $\lambda_k$

### Observations:

- $\lambda$  increases when constraint violated, decreases otherwise
- Constraint violation  $\rightarrow 0$  as training progresses

## Results: Dual-SPMA vs NPG-PD

$J_r$  comparison  
SPMA vs NPG-PD  
(Add plot here)

$J_r(\pi_k)$  vs outer iterations


Constraint violation comparison  
(Add plot here)

Constraint violation vs iterations


### Takeaways:

- Both methods eventually satisfy the constraint
- SPMA: larger, geometry-aware steps
- NPG-PD: smoother but requires careful step-size tuning

## Results: Occupancy Heatmaps



Unconstrained



Constrained

### Left: Unconstrained

Policy explores more broadly

### Right: Safety-Constrained

Policy avoids unsafe states (holes)

# Takeaways & Limitations

## The Recipe:

$$\text{Convex MDP} \xrightarrow{\text{Fenchel}} \text{Saddle-Point Game} \xrightarrow{\text{SPMA}} \text{Shaped-Reward RL}$$

## What we built:

- Dual-SPMA loops for entropy-regularized RL and constrained safety
- NPG-PD baseline for fair comparison
- Three occupancy estimators (tabular MC, feature MC, MLE)

## Limitations:

- Experiments only in low-dimensional (tabular) environments
- MLE estimator not yet stress-tested on large continuous tasks
- Hyperparameter sensitivity not fully characterized

# Future Work

## ① Scale up:

- Test on larger CMDPs (continuous states/actions)
- MuJoCo safety tasks

## ② Better estimation:

- Evaluate MLE estimator on high-dimensional tasks
- Compare variance of different estimators

## ③ More objectives:

- Risk-sensitive RL
- Imitation learning via convex MDP

## ④ Theoretical analysis:

- Convergence rates for Dual-SPMA
- Sample complexity comparison with NPG-PD

# Questions?

**Shervin:** Background & Motivation

**Ahmed:** Problem Formulation & Fenchel Duality

**Pegah:** Related Work & Our Method

**Danielle:** Experiments & Conclusion

## References

- ① **Zahavy, T., et al.** (2021). "Reward is Enough for Convex MDPs." *NeurIPS 2021*. arXiv:2108.06389
- ② **Asad, A., et al.** (2024). "Fast Convergence of Softmax Policy Mirror Ascent." *arXiv:2405.09781*
- ③ **Ding, D., et al.** (2020). "Natural Policy Gradient Primal-Dual Method for Constrained Markov Decision Processes." *NeurIPS 2020*
- ④ **Barakat, A., et al.** (2024). "Reinforcement Learning with General Utilities: Simpler Variance Reduction and Large State-Action Space." *ICML 2024*
- ⑤ **Schulman, J., et al.** (2015). "Trust Region Policy Optimization." *ICML 2015*
- ⑥ **Sutton, R. & Barto, A.** (2018). "Reinforcement Learning: An Introduction." *MIT Press*

## Backup: Flow Constraints (Occupancy Polytope)

The set  $\mathcal{D}$  of valid occupancy measures satisfies **Bellman flow constraints**:

For all states  $s$ :

$$\sum_a d(s, a) = (1 - \gamma)\rho(s) + \gamma \sum_{s', a'} P(s|s', a') d(s', a')$$

Also:  $d(s, a) \geq 0$  for all  $(s, a)$

**Interpretation:**

- Flow into state  $s$  = initial distribution + discounted flow from other states
- This is a **convex polytope** in  $\mathbb{R}^{|S| \times |A|}$

## Backup: Entropy-Regularized Conjugate

For entropy-regularized objective:

$$f(d) = -\langle r, d \rangle + \alpha \sum_{s,a} d(s,a) \log d(s,a)$$

$$f^*(y) = \alpha \log \sum_{s,a} \exp \left( \frac{y(s,a) + r(s,a)}{\alpha} \right)$$

$$\nabla f^*(y) = \text{softmax} \left( \frac{y + r}{\alpha} \right)$$

The gradient is a softmax distribution—very convenient for computation!

## Backup: SPMA with Function Approximation

In function approximation, the SPMA projection step becomes:

$$\theta_{t+1} = \arg \min_{\theta} \sum_s d^{\pi_t}(s) \text{KL}(\pi_{t+1/2}(\cdot|s) \| \pi_{\theta}(\cdot|s))$$

This is a **convex optimization problem** (weighted KL minimization).

Equivalent to **softmax classification**:

- Labels: actions from  $\pi_{t+1/2}$
- Weights: state occupancies  $d^{\pi_t}(s)$
- Features: state representations

# Backup: Algorithm Pseudocode

## Dual-SPMA Algorithm:

- ① Initialize dual variable  $y_1 = 0$ , policy  $\pi_1$  randomly
- ② For  $k = 1, 2, \dots, K_{\text{outer}}$ :
  - ① **Policy step:** Run  $K_{\text{inner}}$  SPMA iterations on shaped reward  $r_{y_k} = -y_k$ 
$$\pi_{k+1} = \text{SPMA}(\pi_k, r_{y_k}, K_{\text{inner}})$$
  - ② **Estimate occupancy:** Collect trajectories, compute  $\hat{d}^{\pi_{k+1}}$
  - ③ **Dual step:** Update dual variable

$$y_{k+1} = y_k + \alpha \left( \hat{d}^{\pi_{k+1}} - \nabla f^*(y_k) \right)$$

- ④ Return final policy  $\pi_K$

## Backup: Notation Summary

Symbol	Meaning
$\mathcal{S}, \mathcal{A}$	State and action spaces
$\pi(a s)$	Policy (probability of action $a$ in state $s$ )
$r(s, a)$	Reward function
$c(s, a)$	Cost function (for CMDPs)
$\gamma$	Discount factor
$d_\pi(s, a)$	Occupancy measure under policy $\pi$
$J(\pi)$	Expected return
$f(d)$	Convex objective over occupancies
$f^*(y)$	Fenchel conjugate of $f$
$y$	Dual variable
$\lambda$	Lagrange multiplier (for CMDPs)
$\tau$	Safety threshold
$\eta$	SPMA step size
$\alpha$	Dual step size